Provisional AI Capability Timeline Distribution

A distributional forecast across AI 2027 agent thresholds and AI Futures Model milestones

Dated July 19, 2026

Bottom line

My current median is WSI around 2032. I assign roughly 40% probability by the end of 2031, 55% by the end of 2032, and 80% by the end of 2035. This is close to JJ's 2031 +/- 1 year view, later than Daniel's fastest implied takeoff, and earlier than the uploaded widening-thesis forecast (~2035 WSI).

Milestone	Median	Central interpretation
AC	late 2028	Full coding automation at a frontier AI company
SAR	early 2030	Qualitatively better than any human at AI research
ASI	mid 2031	Broad superintelligence under the AI Futures definition
WSI	2032	Wildly superintelligent system under JJ's continuity term

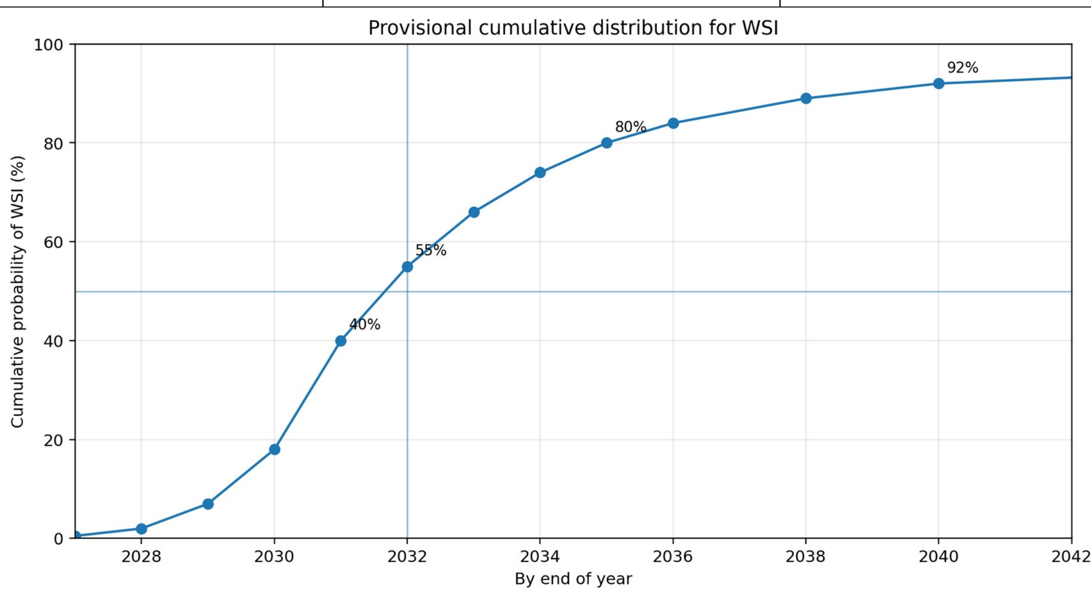

Dates refer to first demonstrated capability at a leading project, not economy-wide deployment. WSI is a continuity term rather than a standardized AI Futures milestone.

1. Distribution table

These quantiles are an explicit, approximate encoding of my current beliefs. They are not a statistical fit and should be revised when definitions or evidence change.

Milestone	p10	p25	p50	p75	p90	Meaning at threshold
Agent 2	2027.0	2027.4	2027.8	2028.2	2029.0	Continual-learning research engineer; about 3x AI R&D uplift
AC	2027.7	2028.1	2028.7	2029.5	2031.2	Company would sooner remove all human software engineers than stop using AI coding labor
SC	2028.0	2028.5	2029.1	2030.0	2032.5	About 900x effective elite coding labor under the AI Futures mapping
SAR	2028.6	2029.2	2030.0	2031.2	2034.0	Qualitatively better than any human at AI research
SIAR	2029.2	2030.0	2031.0	2032.7	2036.5	Research capability gap far beyond SAR
TED-AI	2029.4	2030.2	2031.3	2033.0	2037.0	At least as good as top experts at virtually all cognitive

Milestone	p10	p25	p50	p75	p90	Meaning at threshold
						tasks
ASI	2029.5	2030.3	2031.5	2033.2	2037.5	Broad gap above best humans under the AI Futures definition
WSI	2029.7	2030.7	2032.0	2034.0	2039.0	Wildly superintelligent AI under JJ framework continuity term

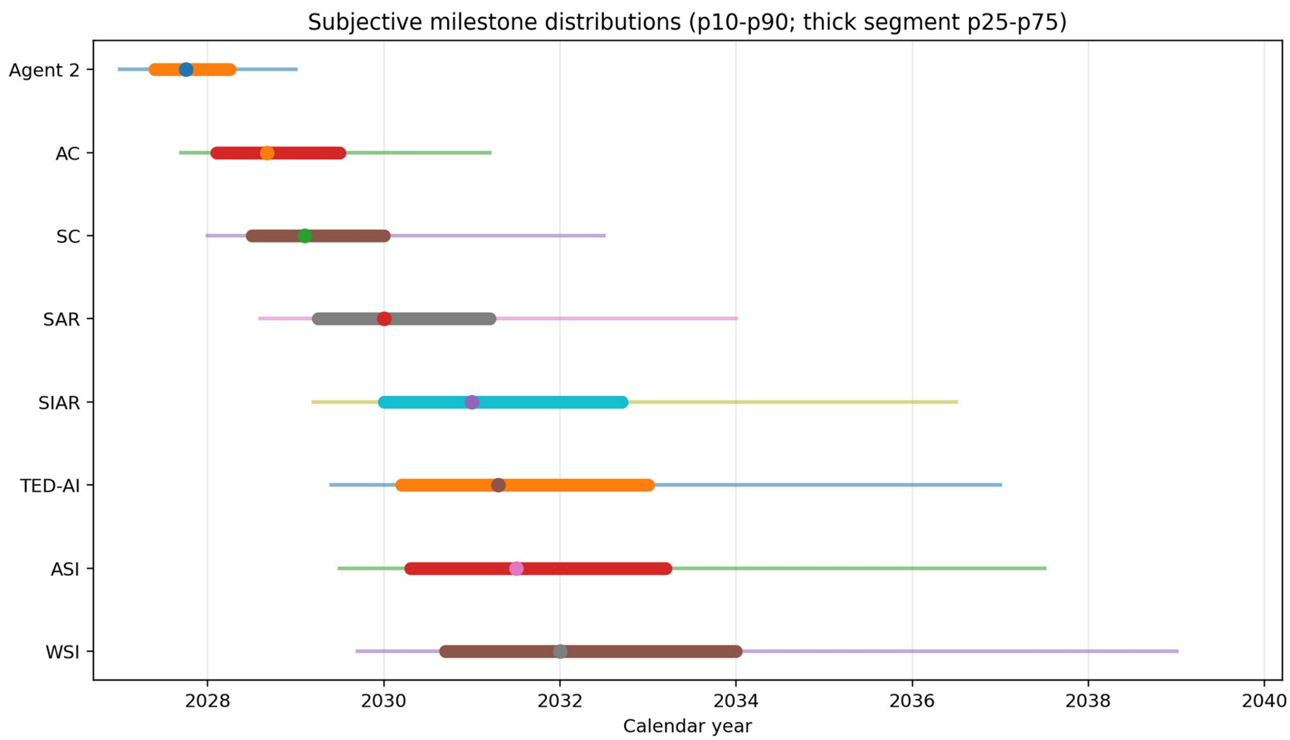

Reading the chart: the dot is the median; the thick segment is p25-p75; the thin line is p10-p90. Wide right tails represent compute, governance, data, evaluation, and research-taste bottlenecks rather than confidence in a single smooth extrapolation.

2. Reasoning summary

Pre-AC: strong evidence for rapid capability growth, but not yet full automation

- Frontier agents are taking on longer and more coherent software tasks, which makes a late-2028 AC median plausible rather than merely a scenario artifact.
- However, benchmark task horizons are not literal uninterrupted work duration. They estimate the amount of serial human labor replaced at a specified success rate, so direct extrapolation to autonomous research should be discounted.
- Continual learning must be disaggregated: persistent memory, frequent offline retraining, online weight updates, self-generated curricula, and improvement of the learning process are different capability thresholds.

AC to SAR: the main widening interval

- Coding automation multiplies experiments, implementation speed, debugging, and parallelism.
- The remaining bottleneck is research taste: choosing high-value questions, noticing mistaken frames, and deciding which negative results should redirect the agenda.
- My median therefore allows about 16 months from AC to SAR - slower than a pure software explosion, but faster than a multi-year diffusion story.

SAR to ASI/WSI: compression resumes

- Once SAR exists, the research-taste bottleneck has partly been crossed by definition. This should accelerate algorithm discovery, evaluation design, synthetic environments, training efficiency, and research coordination.
- Compute and experimental throughput still matter, so I do not put the median on a months-only transition.
- I also do not treat economy-wide diffusion as necessary for first-system WSI. Diffusion can lag capability existence by years.

Core structural claim

The most likely path is neither an immediate AC-to-ASI jump nor a flat six-to-seven-year widening. I expect a meaningful AC-to-SAR delay followed by renewed compression once systems become genuinely superhuman at research taste and agenda selection.

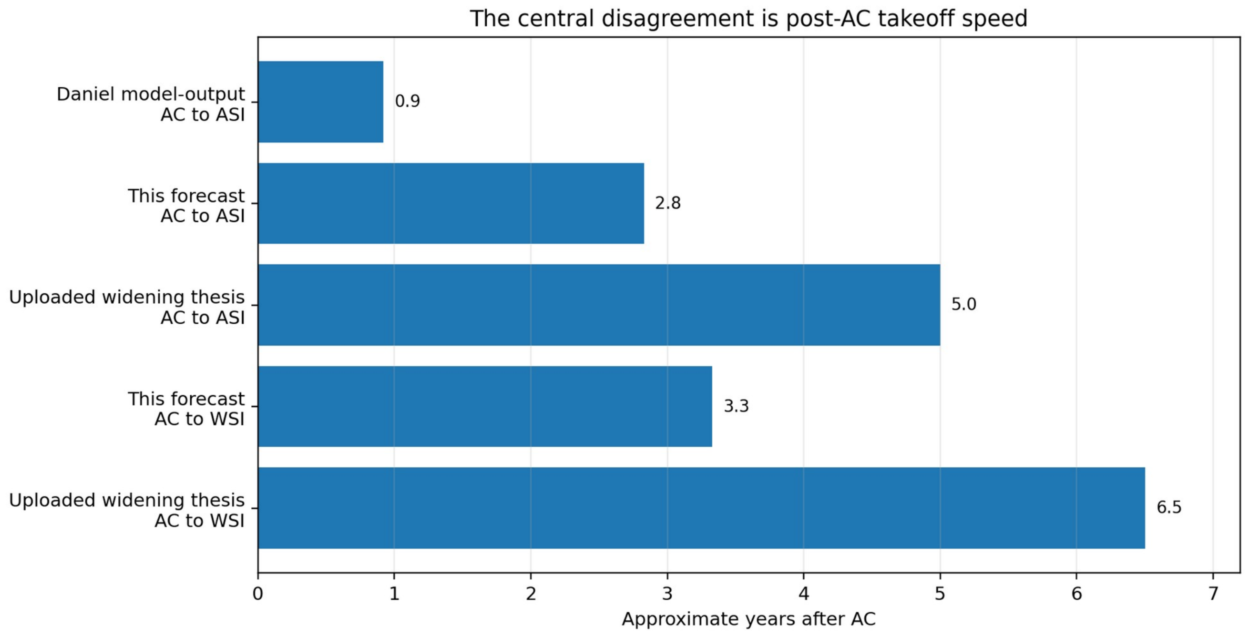
The central disagreement is post-AC takeoff speed
Daniel model-output AC to ASI
This forecast AC to ASI
Uploaded widening thesis AC to ASI
This forecast AC to WSI
Uploaded widening thesis AC to WSI
0.9
2.8
5.0
3.3
6.5
0
1
2
3
4
5
6
7
Approximate years after AC

3. Comparison with JJ, Daniel, and the widening thesis

View	AC median	ASI/WSI implication	Primary disagreement
JJ	~2028, with weight on Daniel	WSI ~2031 +/- 1 year	Moderately fast post-AC transition
Daniel (Q1 2026/model output)	06/2028	ASI 05/2029 model output	Very fast AC-to-ASI takeoff
This forecast	late 2028	ASI mid-2031; WSI 2032	Widening to SAR, then compression
Uploaded widening thesis	mid-2028	ASI mid-2033; WSI ~2035	Long post-AC bottlenecks and diffusion

Agreement with Daniel: The late-2028 AC estimate is credible. The frontier evidence is now consistent with rapidly improving autonomous coding and research engineering.

Disagreement with Daniel: An approximately eleven-month AC-to-ASI median places too little weight on research taste, experimental compute, coordination, and the difficulty of validating new paradigms.

Agreement with the widening thesis: Automating perspiration does not automatically automate agenda selection. The AC-to-SAR interval is where widening is most plausible.

Disagreement with the widening thesis: Once SAR or SIAR exists, a multi-year delay to first-system WSI seems too slow unless compute or governance constraints become binding. Slow diffusion should not be conflated with slow capability formation.

Why JJ's 2031 +/- 1 remains defensible

It implies roughly 2.5 to 3.5 years from AC to WSI. That gives genuine weight to bottlenecks while still reflecting a positive-feedback process after autonomous superhuman AI research appears. My one-year-later median is caution, not a categorical structural disagreement.

4. Falsifiers and forecast-update rules

The forecast should move only when evidence bears on a registered mechanism. Each trigger should define the observation, confirmation standard, and directional update before the result is known.

Trigger	Pre-registered observation	Forecast implication
Continual learning - operational	A deployed frontier system acquires durable new competence from ongoing experience, with weight-level or equivalent adaptation, measured retention, bounded interference, and generalization beyond memorized episodes.	Earlier Agent-2 and AC; modestly shorter AC-to-SAR if the learning loop is autonomous.
Novel direction in a hard-to-verify domain	Given a broad objective rather than a supplied agenda, the system originates a non-obvious research direction, designs discriminating experiments, and produces a result that blind domain experts judge novel and valid.	Strongly shorter AC-to-SAR and SAR-to-ASI; major evidence against the widening thesis.
Autonomous AI R&D loop	The system proposes, implements, evaluates, interprets, and redirects a sequence of AI experiments with limited human choice points, producing a measured algorithmic gain on held-out evaluations.	Moves SAR and downstream milestones earlier; effect depends on compute consumed and human oversight.
Compute-bound plateau	Large increases in agent labor fail to raise algorithmic progress because decisive experiments are limited by training or inference compute, hardware delivery, power, or validation latency.	Lengthens SAR-to-ASI/WSI; supports widening after SAR.
Broad expert dominance before SAR	Systems dominate top humans across messy, slow-feedback domains while remaining only SAR-level in AI research.	Evidence against the ordering SIAR before TED-AI and against research-domain-first compression.

5. Sources, definitions, and caveats

Confirmation and update discipline

Prefer blind expert assessment, public artifacts, and at least one independent replication or equivalent internal audit. Check for benchmark leakage, tool assistance, and hidden human intervention. For each revision, record the prior distribution, evidence card, mechanism affected, directional update, new distribution, reviewer, and timestamp; never silently overwrite earlier forecasts.

Primary definition sources

- AI Futures Project - AI Futures Model: Dec 2025 Update
- AI Futures Project - Q1 2026 Timelines Update
- AI 2027 scenario and research supplements
- METR - Task-Completion Time Horizons of Frontier AI Models
- Uploaded: Claude Widening Thesis Timeline Jul2026.pdf

Important caveats

- This is a subjective forecast and a compact reasoning summary, not a claim of mathematically objective probabilities.
- Definitions must be versioned. AC, SC, SAR, SIAR, TED-AI, ASI, and AI 2027 Agent labels are related but not interchangeable.
- The far-right tail is especially uncertain. Governance interventions, accidents, wars, supply constraints, paradigm stagnation, or deliberate pauses can dominate capability extrapolation.
- Medians can obscure branch structure. A fast software-intelligence-explosion branch and a slower compute- or taste-limited branch can produce the same median.
- Evidence from laboratory evaluations may not generalize to open-ended scientific work; conversely, private frontier-lab capabilities may lead public benchmarks.

Recommended revisit cadence

Conduct a structured review twice per month for evidence intake, but revise the public distribution only when a registered trigger is crossed or a cluster of weaker evidence produces a material cumulative update. Publish a full re-derivation at least quarterly.

Prepared as a decision aid for JJ's forecasting framework. The purpose is to make disagreements legible and falsifiable, not to create a false impression of precision.